

# **polars-bio: High-Performance Python DataFrame Operations for Genomics**

Demystify AI & Data Management Series

Marek Wiewiórka

September 15, 2025

## About me



- ▶ Assistant Professor<sup>a</sup> at Warsaw University of Technology
- ▶ Chief Architect @Xebia Data Poland, 20+ years building data-intensive systems
- ▶ distributed and data-intensive systems, artificial intelligence and cloud computing for large scale genomic studies.
- ▶ road and gravel bikes enthusiast
- ▶ <https://marekwiewiorka.org/>

---

<sup>a</sup>Institute of Computer Science

# Biodatageeks lab

- Warsaw University of Technology,  
Faculty of Electronics and  
Information Technology
-  current research topics:
  - AI for analyzing biomedical literature
  - Meta-calling for gene fusion detection in RNA-Seq
  - Optimizing RVAS
  - Open genomic data lakehouse
- <https://biodatageeks.org/>
- <https://github.com/biodatageeks/>

## Meet the Team

### Principal Investigators



Tomasz Gambin  
Associate Professor



Marek Wiewiórka  
Assistant Professor

### Researchers



Anna Kosycarz  
BSc Student



Iga Ostrowska  
PhD student



Wojciech Sitek  
Research Assistant



Piotr Suszyński  
PhD student



Agnieszka Szmułko  
PhD Student

# Agenda

1. 🤔 Rationale and motivation
2. 🌎 Context and alternatives for polars-bio
3. 🔬 Deep dive into internals
4. 📊 Benchmarks
5. 🌟 Future directions

# Introduction to polars-bio

- ▶ polars-bio is a novel Python DataFrame library for genomics that is *fast* and *memory-efficient*, introduced in 2025, built on top of Polars, Apache DataFusion and Apache Arrow.
- ▶ main focus areas:
  - ▶ 🧪 genomic interval operations
  - ▶ 🚀 scalable data processing and querying
  - ▶ 💾 fast I/O for bioinformatics file formats
  - ▶ ☁ cloud storage interoperability
  - ▶ 🏢 genomic data lakehouse readiness



POLARS-BIO



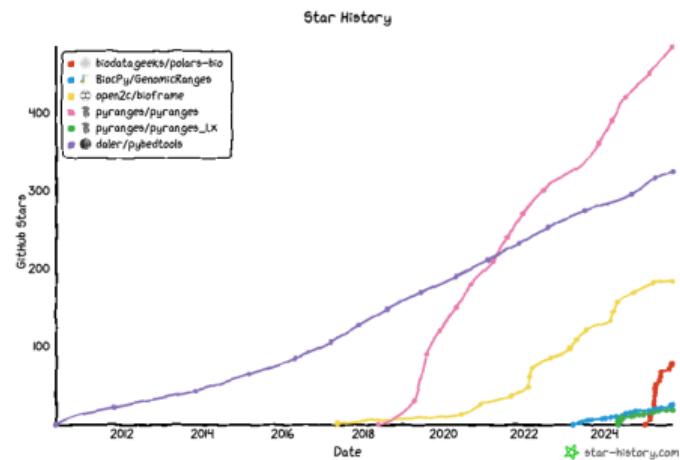
## Rationale, History, and Challenges

- ▶  Growing bioinfo dataset sizes vs. increasing capacity of commodity hardware
- ▶  Trade-off: scalability of distributed systems (e.g., Apache Spark - Hail, Glow) vs. simplicity and performance of single-node libraries (e.g. DuckDB)
- ▶  Single-node solutions: constrained in both performance and scalability
- ▶  First attempt (2019–2023): SeQuiLa project on top of Apache Spark
- ▶  Conclusion: towards a *hybrid* approach



# Landscape of tools for genomic interval operations in Python

- ▶ several widely used libraries exist in this space:
  - ▶ Pyranges and new Pyranges1
  - ▶ Pybedtools
  - ▶ Bioframe
  - ▶ GenomicRanges
- ▶ employing an *eager, in-memory* execution model with Pandas DataFrames/ NumPy arrays
- ▶ sweep-line (Bioframe, Pyranges1) or Nested Containment List (Pyranges, GenomicRanges) or genome binning algorithm (Pybedtools)
- ▶ focus *primarily* on optimizing genomic operations rather than end-to-end processing and IO operations



## Market trends in data systems

- ▶ out-of-core (streaming) processing
- ▶ single node vectorized engines – e.g. DuckDB, Polars
- ▶ *lazy* evaluation and query optimization – e.g. Polars
- ▶ open data standards and interoperability,  
such as Apache Arrow or Apache Iceberg
- ▶ composable and reusable, e.g. query parsers,  
optimizers, query engines, memory and file/table formats
- ▶ data lakehouse architecture



DuckDB

## Composable Data Management Systems (CDMS) Manifesto

- ▶ **Problem:** Data systems are *fragmented, duplicated*, hard to maintain
- ▶ **Vision:** Break *monoliths* into *modular, reusable* components (frontends, Internal Representation, optimizers, execution engines, runtime environments)
- ▶ **Why Now:** Already existing *open standards* (Arrow, Parquet, Iceberg) enable composability
- ▶ **Examples:** Velox, Apache DataFusion
- ▶ **Benefits:** *Faster* innovation, *reduced* engineering effort, consistent user experience

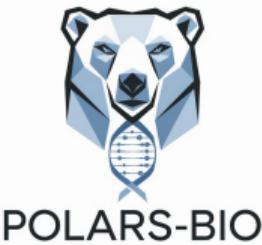
## Limitations of Current Approaches to Genomic Interval Processing

Genomic intervals processing is closer to BI/DWH/ETL-style workloads than to numerical computing!

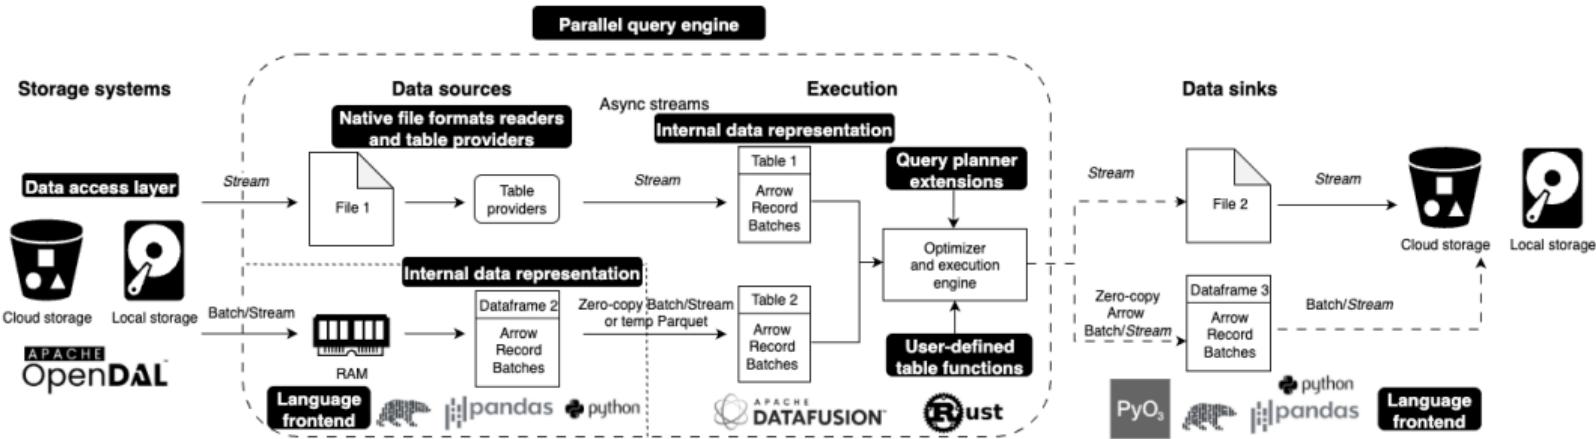
- ▶ Relying on libraries (e.g., NumPy) not designed for efficient bioinformatics data handling
- ▶ Re-implementing algorithms and reinventing the wheel instead of leveraging mature *query engine*: optimizers, operators and open data standards
- ▶ *Parallelism* and *out-of-core* not treated as a first-class concern (limited scalability)
- ▶ *Naive Python* implementations (slow, limited scalability)
- ▶ Missing *end-to-end optimization* including reading, processing and writing data

## Why polars-bio is different?

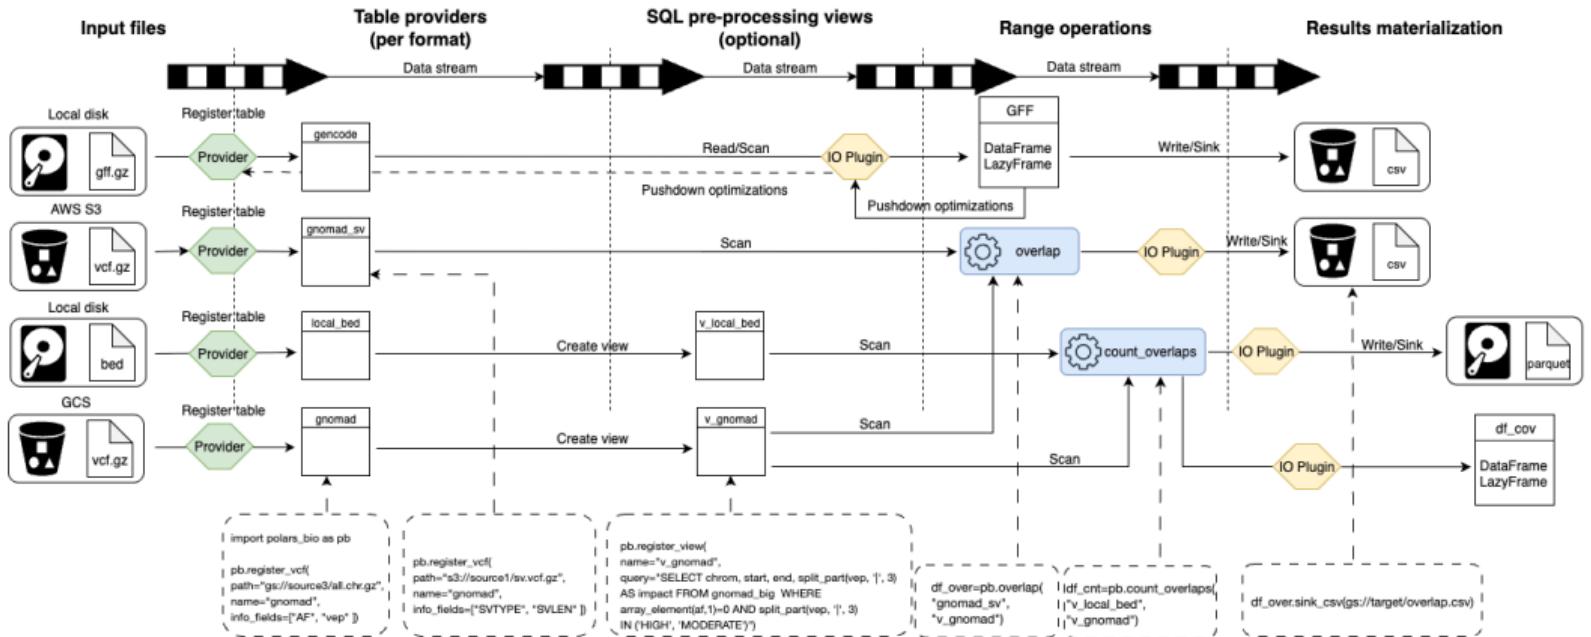
- ▶ Composable and best of breed approach
  - ▶ query engine (Apache DataFusion)
  - ▶ DataFrame library (Polars)
  - ▶ columnar memory format (Apache Arrow)
  - ▶ data structure for interval intersection queries (COITrees and Superintervals)
  - ▶ bioinformatics file formats (noodles)
- ▶ Builtin *lazy*, *out-of-core* and *parallel* computational model
- ▶ IO layer optimizations for analytical queries, such as *projection* and *predicate* pushdowns



# polars-bio high-level architecture



# Architecture deep-dive - core components



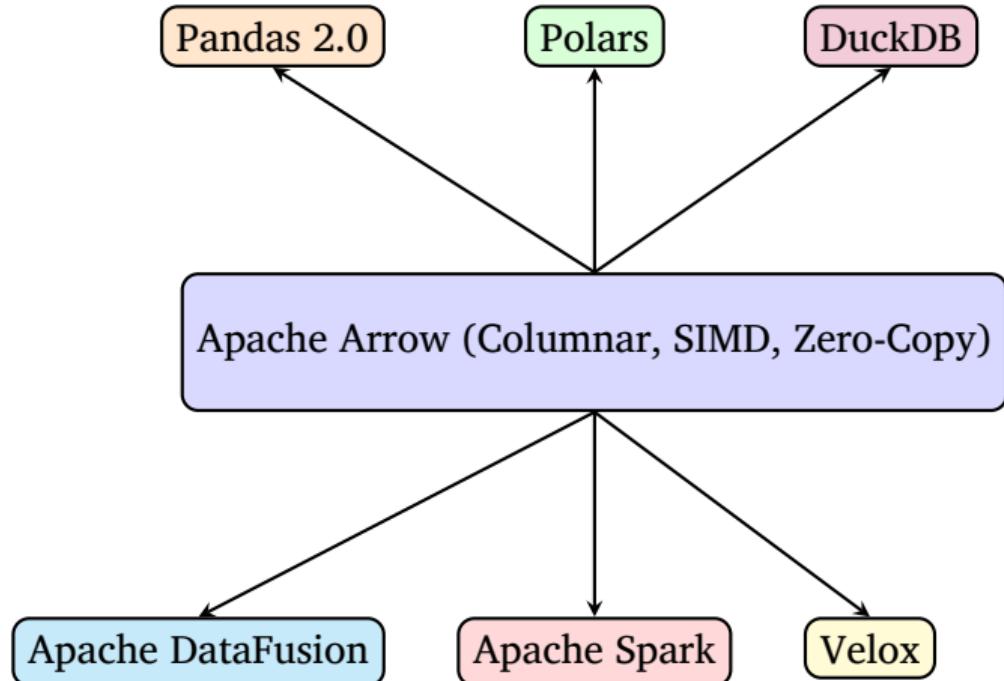
## Architecture deep-dive - Polars and Apache DataFusion primer

- ▶ Polars and Apache DataFusion exhibit significant similarities, such as Apache Arrow columnar memory model, lazy evaluation and out-of-core computational model, great performance
- ▶ different main focuses:
  - ▶ Polars – feature-rich end-user DataFrame library
  - ▶ DataFusion – extremely extensible query engine for building custom data systems
- ▶ do we really need both?
  - ▶ Polars' great data wrangling capabilities but hard to extend
  - ▶ DataFusion's codebase reusability (e.g. hybrid execution) and more robust abstractions for query and IO optimizations
  - ▶ additional integration complexity (e.g. pushdown optimizations, parallelism control)



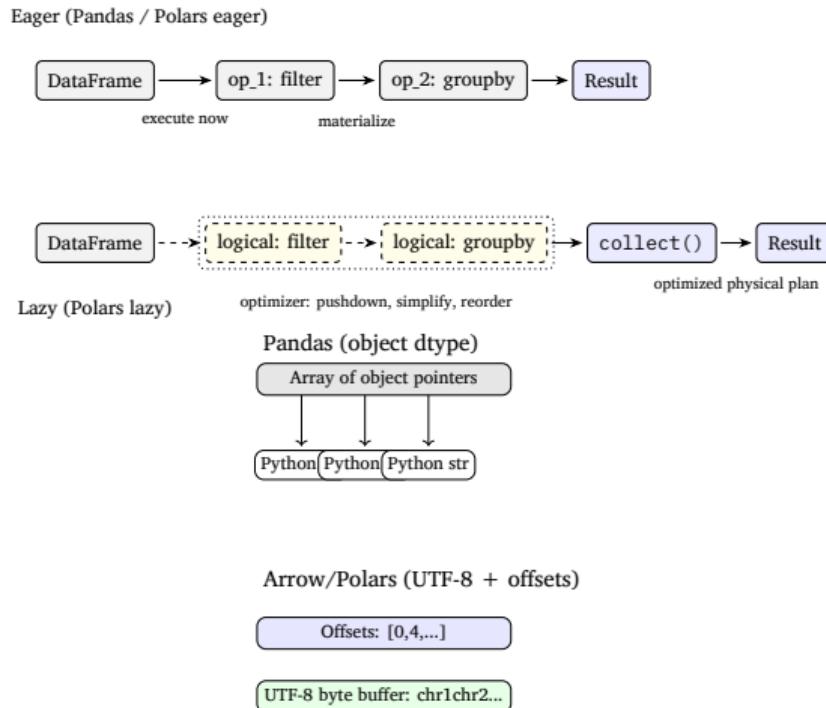
# Architecture deep-dive - Apache Arrow

- ▶ Standardized columnar memory format – zero-copy sharing
- ▶ Vectorized execution: SIMD and CPU cache efficiency
- ▶ Cross-language interoperability (e.g. Python and Rust)
- ▶ Integration with open standards – Parquet, Iceberg
- ▶ Foundation for modern data systems – Polars, Ray, Rapids, Apache Spark



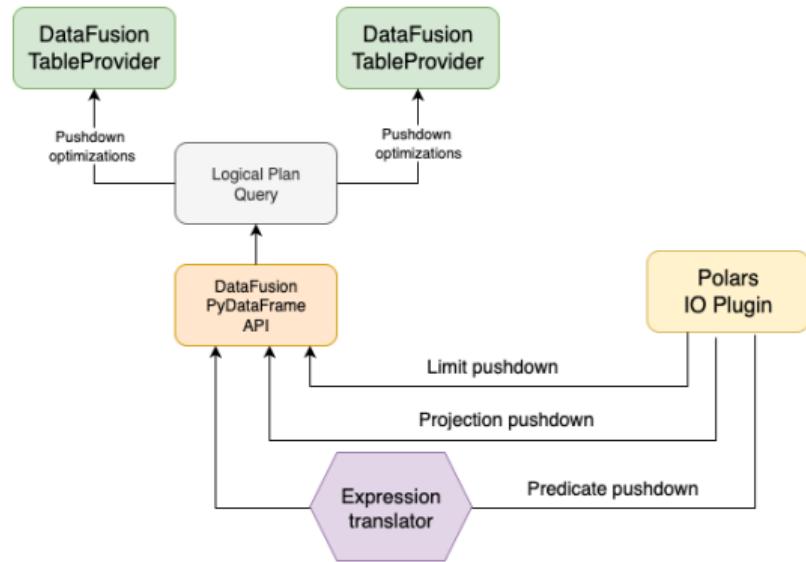
# Architecture deep-dive - Polars vs. Pandas

- ▶ Execution model: Pandas is eager-only; Polars supports eager *and* lazy.
- ▶ Optimization: Pandas has no query optimizer; Polars (lazy) performs projection/predicate pushdown, simplification, reordering.
- ▶ Parallelism: Pandas mostly single-threaded (Python/GIL); Polars is multi-threaded (Rust).
- ▶ Memory/layout: Pandas uses NumPy blocks; Polars is columnar and Arrow-friendly.
- ▶ Out-of-core/streaming: Pandas primarily in-memory; Polars supports streaming/out-of-core in lazy plans.
- ▶ String handling: Pandas often stores Python objects (high memory overheads); Polars stores UTF-8 natively with efficient kernels (SIMD).



# Architecture deep-dive - Polars IO plugin

- ▶ arbitrary function that returns a generator (Iterator) producing `pl.DataFrame` batches and gets back `LazyFrame`
- ▶ used for both files scanning and interval operations results streaming
- ▶ zero-copy and streaming using Arrow RecordBatchStream with DataFusion PyDataFrame
- ▶ support for limit, projection and predicate pushdowns (currently only GFF)



## Architecture deep-dive - input file formats

- ▶ subproject datafusion-bio-formats
- ▶ exposed using custom TableProviders
- ▶ support for parallel reading of BGZF inputs
- ▶ local and cloud storage (AWS S3, GCS and Azure Blob)
- ▶ cloud storage supported features

Format	Single-threaded	Parallel	Limit pushdown	Predicate pushdown	Projection pushdown
BED	✓	✗	✓	✗	✗
VCF	✓	✗	✓	✗	✗
BAM	✓	✗	✓	✗	✗
FASTQ	✓	✓	✓	✗	✗
FASTA	✓	✗	✓	✗	✗
GFF3	✓	✓	✓	✓	✓

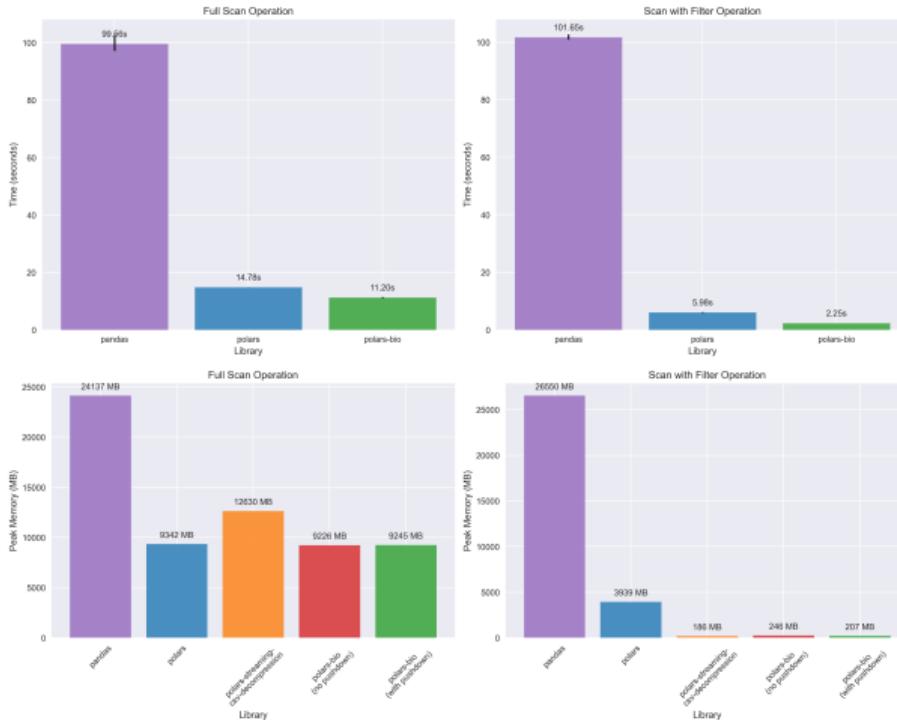
## Benchmarking dataset

- ▶ AIList real dataset converted into Parquet format – details
- ▶ GFF3 GENCODE release 49

Dataset#	Name	Size(x1000)	Description
0	chainRn4	2,351	<a href="#">Source</a>
1	fBrain	199	<a href="#">Source</a>
2	exons	439	Dataset used in the BEDTools tutorial.
3	chainOrnAna1	1,957	<a href="#">Source</a>
4	chainVicPac2	7,684	<a href="#">Source</a>
5	chainXenTro3Link	50,981	<a href="#">Source</a>
6	chainMonDom5Link	128,187	<a href="#">Source</a>
7	ex-anno	1,194	Dataset contains GenCode annotations with ~1.2 million lines, mixing all types of features.
8	ex-rna	9,945	Dataset contains ~10 million direct-RNA mappings.

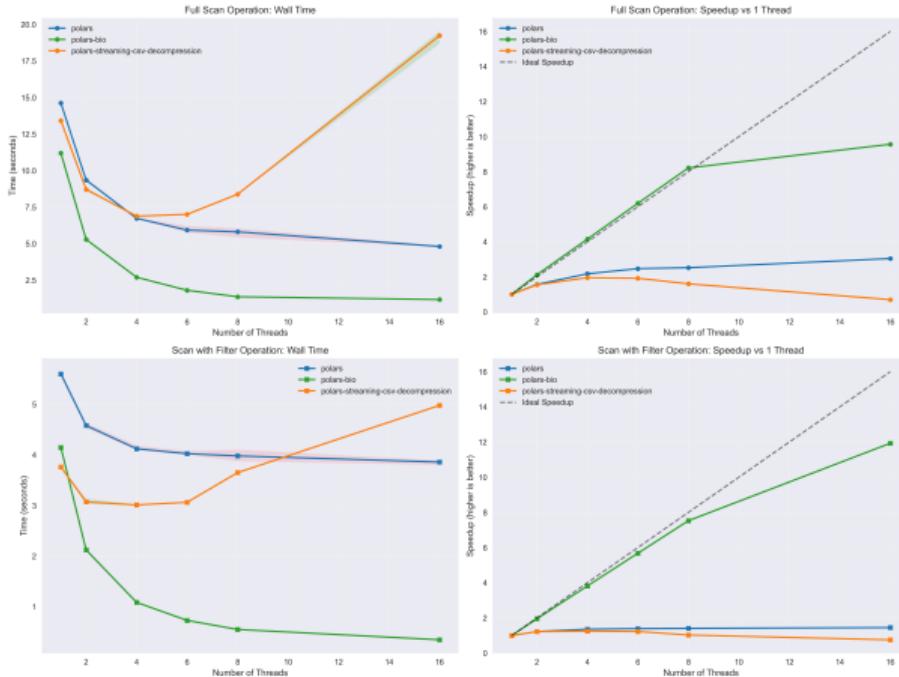
Source: Jianglin Feng , Aakrosh Ratan , Nathan C Sheffield, *Augmented Interval List: a novel data structure for efficient genomic interval search*, Bioinformatics 2019.

# File formats: GFF (scan\_csv vs polars-bio) – results 1/2



- ▶ in full-scans Polars and polars-bio significantly outperform Pandas
- ▶ Polars problem with scan\_csv and compressed files)
- ▶ streaming decompression plugin

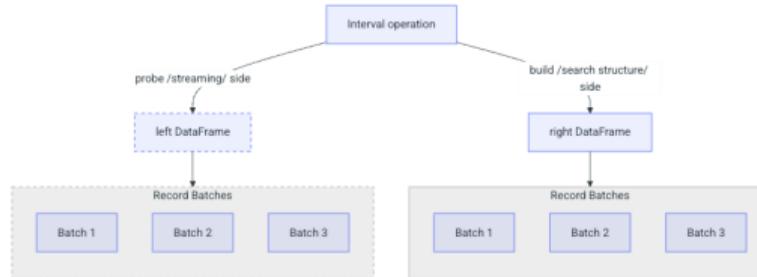
## File formats: GFF (scan\_csv vs polars-bio) – results 2/2



- polars-bio achieves near-linear scaling up to 8 threads
- Polars and streaming decompression plugin scale poorly

# Architecture deep-dive – genomic interval operations 1/2

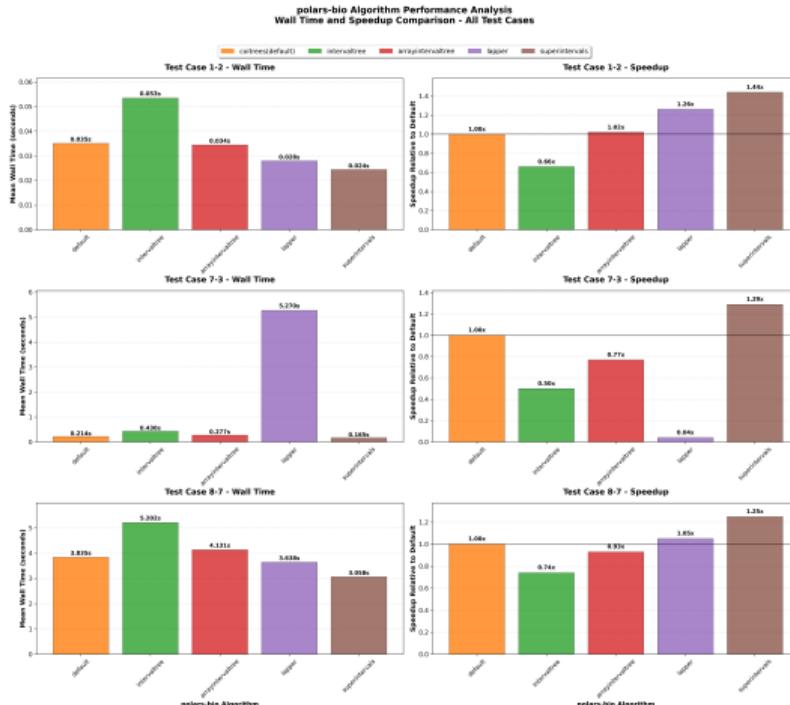
- ▶ inspired by the Hash Join implementation in DataFusion
- ▶ the *entire* (coordinates) build side is read into the interval search data structure
- ▶ batches from the probe side are *streamed* through and checked against the contents of the search data structure



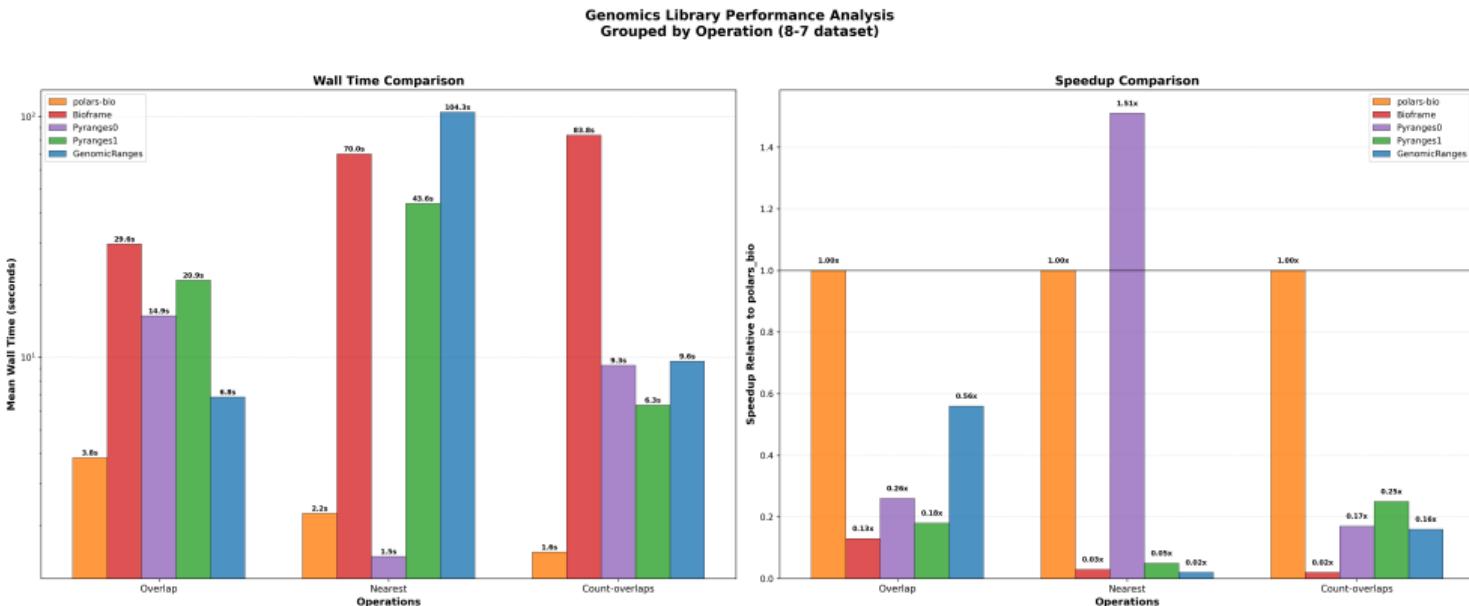
- ▶ subproject sequila-native
- ▶ custom PhysicalPlanner and PhysicalOptimizerRule for detecting and rewriting generic interval join operation (overlap or nearest)
- ▶ User-Defined Table Function (UDTF) for operations, such as coverage or count overlaps
- ▶ several data structures available:
  - ▶ COITrees
  - ▶ IITree
  - ▶ AVL-tree
  - ▶ rust-lapper
  - ▶ Superintervals

# Genomic interval operations – structures comparison results 1/6

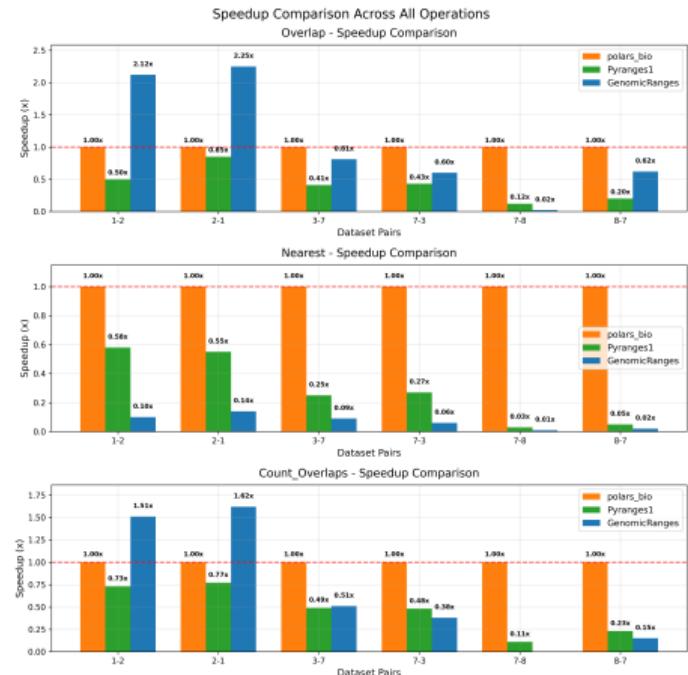
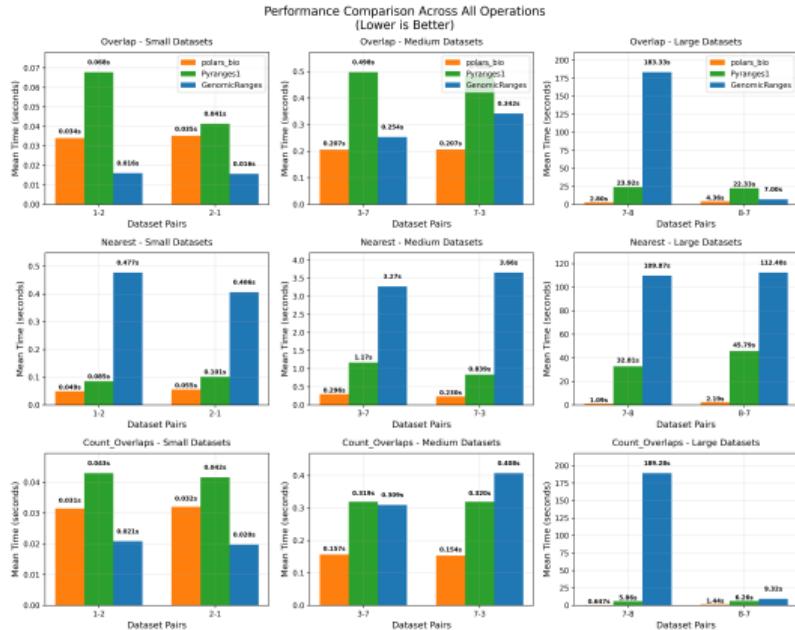
- ▶ COITrees  
(polars-bio default)  
and Superintervals  
fastest in all test  
cases
- ▶ configurable in  
runtime
- ▶ more tests using  
different datasets  
characteristics  
needed



# Genomic interval operations – results 2/6

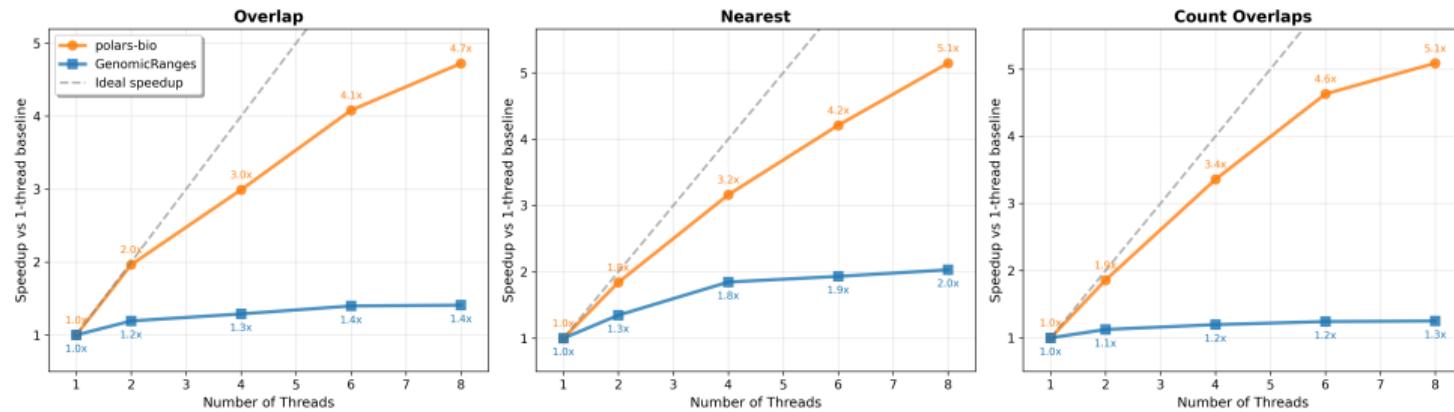


# Genomic interval operations – results 3/6

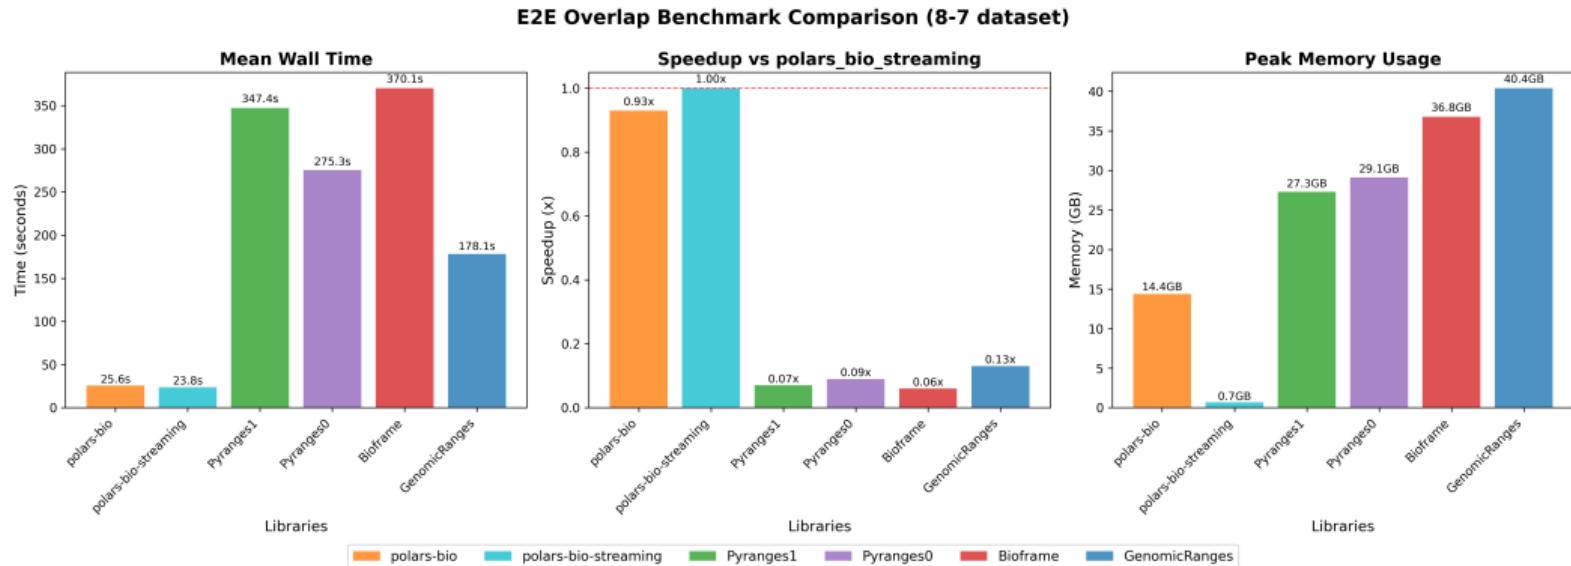


## Genomic interval operations – scaling – results 4/6

Parallel Scaling Performance (8-7 dataset)  
Speedup relative to 1-thread baseline



## Genomic interval operations – e2e pipeline – results 5/6



## Genomic interval operations – results 6/6

```
import pandas as pd
>> df = pd.read_parquet("/tmp/exons/")

# GenomicRanges - int32
GenomicRanges(number_of_ranges=438694, seqnames=[], ranges=IRanges(
    start=array([], shape=(438694,), dtype=int32),
    width=array([], shape=(438694,), dtype=int32)),

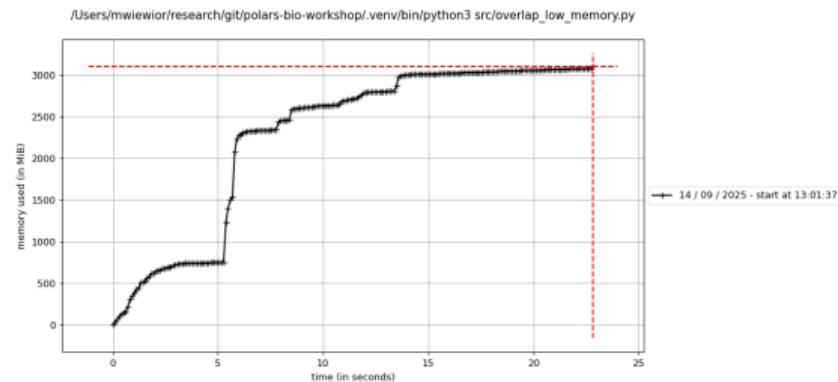
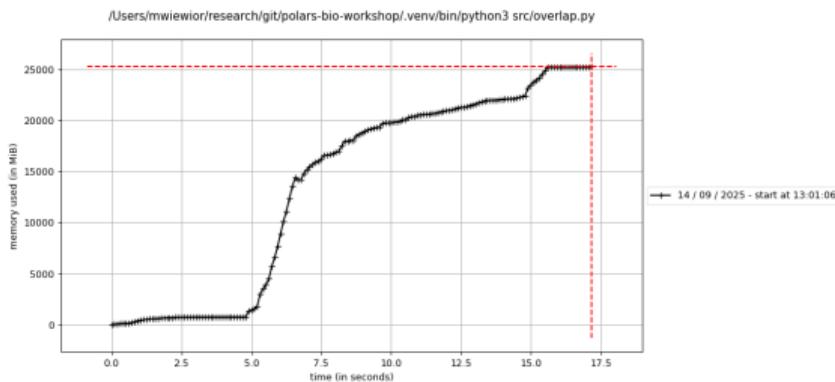
# Bioframe - int32
cols=["contig","pos_start","pos_end"]
>> bf.from_any(df, cols=cols).info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 438694 entries, 0 to 438693
Data columns (total 3 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   contig      438694 non-null  object  
 1   pos_start   438694 non-null  int32  
 2   pos_end     438694 non-null  int32  
dtypes: int32(2), object(1)
```

```
# Pyranges0 - int64 !!!
>> df2pr0(df)
+-----+-----+-----+
| Chromosome | Start | End |
| (category) | (int64) | (int64) |
+-----+-----+-----+
| chr1       | 11873 | 12227 |
| chr1       | 12612 | 12721 |
| chr1       | 13220 | 14409 |
| chr1       | 14361 | 14829 |
| ...        | ...   | ...  |
+-----+-----+-----+
Unstranded PyRanges object has 438,694 rows and 3 columns.

# Pyranges1 - int32
>> df2pr1(df)
index | Chromosome | Start | End
int64 | object     | int32 | int32
+-----+-----+-----+
0     | chr1       | 11873 | 12227
1     | chr1       | 12612 | 12721
2     | chr1       | 13220 | 14409
3     | chr1       | 14361 | 14829
...
PyRanges with 438694 rows, 3 columns, and 1 index columns.
```

## Overlap operation low memory mode for $\sim 10^9$ wide rows

Capped (max rows per batch) streaming-friendly emission –  $\sim 30 - 50\%$  slower but with significantly lower memory utilization.



- ▶  Lakehouse support with open standards
  - ▶  Apache Iceberg integration with open-source (e.g. Apache Gravitino, Lakekeeper, Apache Polaris, Unity Catalog OSS) and proprietary catalogs
- ▶  Feature parity across all supported bioinformatics formats
- ▶  Write-back into table formats (e.g. Apache Iceberg)
- ▶  Spec-driven agentic development for automated pipelines
- ▶  Hybrid execution: SeQuila + Apache Comet accelerator
- ▶  **Your use case!**

- ▶  polars-bio: a new Python DataFrame library for genomics
- ▶  Combines Polars, Apache DataFusion, and Apache Arrow for speed and scalability
- ▶  Efficient I/O for popular bioinformatics formats
- ▶  Addresses limitations of existing interval processing tools
- ▶  Towards a hybrid, lakehouse-ready approach for large-scale genomics

Stay Tuned!

- ▶  GitHub: [github.com/biodatageeks/polars-bio](https://github.com/biodatageeks/polars-bio)
- ▶  Project page: [biodatageeks.org/polars-bio](https://biodatageeks.org/polars-bio)
- ▶  Discord: Join our community
- ▶  Meet us at **ASHG 2025 Annual Meeting, Boston, October 14-18, 2025**

# Thank You!



## Questions

## Hands-on Demo



[github.com/biodatageeks/polars-bio-workshop](https://github.com/biodatageeks/polars-bio-workshop)